

1 Data Structures

pandas introduces two new data structures to Python - Series and DataFrame, both of which are built on top of NumPy (this means it's fast).

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
pd.set_option('max_columns', 50)
```

1.1 Series

Series is a one-dimensional labeled array capable of holding any data type (integers, strings, floating point numbers, Python objects, etc.). The axis labels are collectively referred to as the index. The basic method to create a Series is to call:

```
s = Series(data, index=index)
```

Here, data can be many different things:

- a Python dict
- an ndarray
- a scalar value (like 5)

A Series is a one-dimensional object similar to an array, list, or column in a table. It will assign a labeled index to each item in the Series. By default, each item will receive an index label from 0 to N, where N is the length of the Series minus one.

```
# create a Series with an arbitrary list
s = pd.Series([7, 'Heisenberg', 3.14, -1789710578, 'Happy Eating!'])
s
```

```
0          7
1    Heisenberg
2         3.14
```

Data Analysis with Python

```
3      -1789710578
4      Happy Eating!
dtype: object
```

Alternatively, you can specify an index to use when creating the Series.

```
s = pd.Series([7, 'Heisenberg', 3.14, -1789710578, 'Happy Eating!'],
              index=['A', 'Z', 'C', 'Y', 'E'])
s
```

```
A      7
Z      Heisenberg
C      3.14
Y      -1789710578
E      Happy Eating!
dtype: object
```

The Series constructor can convert a dictionary as well, using the keys of the dictionary as its index.

```
d = {'Chicago': 1000, 'New York': 1300, 'Portland': 900, 'San Francisco': 1100,
      'Austin': 450, 'Boston': None}
cities = pd.Series(d)
cities
Out[4]:
Austin      450
Boston      NaN
Chicago     1000
New York    1300
Portland     900
San Francisco 1100
dtype: float64
```

You can use the index to select specific items from the Series ...

```
cities['Chicago']
Out[5]:
1000.0
```

```
cities[['Chicago', 'Portland', 'San Francisco']]
Out[6]:
Chicago          1000
Portland          900
San Francisco    1100
dtype: float64
```

Or you can use boolean indexing for selection.

```
cities[cities < 1000]
Out[7]:
Austin          450
Portland        900
dtype: float64
```

That last one might be a little weird, so let's make it more clear - `cities < 1000` returns a Series of True/False values, which we then pass to our Series `cities`, returning the corresponding True items.

```
less_than_1000 = cities < 1000
print less_than_1000
print '\n'
print cities[less_than_1000]
Austin          True
Boston          False
Chicago         False
New York        False
Portland         True
San Francisco   False
dtype: bool

Austin          450
Portland        900
dtype: float64
```

You can also change the values in a Series on the fly.

```
\begin{framed}
\begin{verbatim}
# changing based on the index
print 'Old value:', cities['Chicago']
cities['Chicago'] = 1400
print 'New value:', cities['Chicago']
Old value: 1000.0
New value: 1400.0

\begin{framed}
\begin{verbatim}
# changing values using boolean logic
print cities[cities < 1000]
print '\n'
cities[cities < 1000] = 750

print cities[cities < 1000]
Austin      450
Portland    900
dtype: float64

Austin      750
Portland    750
dtype: float64
```

What if you aren't sure whether an item is in the Series? You can check using idiomatic Python.

```
print 'Seattle' in cities
print 'San Francisco' in cities
False
True
```

Mathematical operations can be done using scalars and functions.

```
# divide city values by 3
cities / 3
Out[12]:
Austin          250.000000
Boston          NaN
Chicago         466.666667
New York        433.333333
Portland        250.000000
San Francisco   366.666667
dtype: float64
```

```
\begin{framed}
\begin{verbatim}
# square city values
np.square(cities)
Out[13]:
Austin          562500
Boston          NaN
Chicago        1960000
New York        1690000
Portland        562500
San Francisco   1210000
dtype: float64
```

You can add two Series together, which returns a union of the two Series with the addition occurring on the shared index values. Values on either Series that did not have a shared index will produce a NULL/NaN (not a number).

```
print cities[['Chicago', 'New York', 'Portland']]
print'\n'
print cities[['Austin', 'New York']]
print'\n'
print cities[['Chicago', 'New York', 'Portland']] + cities[['Austin', 'New York']]
```

```
Chicago        1400
New York        1300
Portland        750
dtype: float64
```

```
Austin      750
New York    1300
dtype: float64
```

```
Austin      NaN
Chicago     NaN
New York    2600
Portland    NaN
dtype: float64
```

Notice that because Austin, Chicago, and Portland were not found in both Series, they were returned with NULL/NaN values.

NULL checking can be performed with `isnull` and `notnull`.

```
# returns a boolean series indicating which values aren't NULL
cities.notnull()
Out[15]:
Austin      True
Boston      False
Chicago     True
New York    True
Portland    True
San Francisco True
dtype: bool
In [16]:
# use boolean logic to grab the NULL cities
print cities.isnull()
print '\n'
print cities[cities.isnull()]
Austin      False
Boston      True
Chicago     False
New York    False
Portland    False
San Francisco False
dtype: bool

Boston      NaN
```

Data Analysis with Python

```
dtype: float64
```

1.2 DataFrame

A DataFrame is a tabular data structure comprised of rows and columns, akin to a spreadsheet, database table, or R's `data.frame` object. You can also think of a DataFrame as a group of Series objects that share an index (the column names).

For the rest of the tutorial, we'll be primarily working with DataFrames.

1.3 Panel

Panel is a somewhat less-used, but still important container for 3-dimensional data. The term panel data is derived from econometrics and is partially responsible for the name pandas: pan(el)-da(ta)-s. The names for the 3 axes are intended to give some semantic meaning to describing operations involving panel data and, in particular, econometric analysis of panel data. However, for the strict purposes of slicing and dicing a collection of DataFrame objects, you may find the axis names slightly arbitrary:

- `items`: axis 0, each item corresponds to a DataFrame contained inside
- `major_axis`: axis 1, it is the index (rows) of each of the DataFrames
- `minor_axis`: axis 2, it is the columns of each of the DataFrames