

1 Basic Functions and Numerical Indexing

1.1 Mathematics

sum, cumsum

sum sums elements in an array. By default, it will sum all elements in the array, and so the second argument is normally used to provide the axis to use – 0 to sum down columns, 1 to sum across rows. cumsum produces the cumulative sum of the values in the array, and is also usually used with the second argument to indicate the axis to use.

```
>>> x = randn(3,4)
>>> x
array([[0.08542071,
 2.05598312,
 2.1114733 , 0.7986635 ],
 [0.17576066,
 0.83327885, 0.64064119,
 0.25631728],
 [0.38226593,
 1.09519101,
 0.29416551, 0.03059909]])
>>> sum(x) # all elements
0.62339964288008698
>>> sum(x, 0) # Down rows, 4 elements
array([0.6434473
 , 2.31789529,
 1.76499762, 0.57294532])
>>> sum(x, 1) # Across columns, 3 elements
array([ 0.76873297, 0.23944028,
 1.15269233])
>>> cumsum(x,0) # Down rows
array([[0.08542071,
 2.05598312,
 2.1114733 , 0.7986635 ],
 [0.26118137,
 1.22270427,
 1.47083211, 0.54234622],
 [0.6434473
 , 2.31789529,
 1.76499762, 0.57294532]])
```

sum and cumsum can both be used as function or as methods. When used as methods, the first input is the axis so that `sum(x,0)` is the same as `x.sum(0)`.

prod, cumprod

prod and cumprod behave similarly to sum and cumsum except that the product and cumulative product are returned. prod and cumprod can be called as function or methods.

diff

diff computes the finite difference of a vector (also array) and returns an $n-1$ element vector when used on an n element vector. diff operates on the last axis by default, and so `diff(x)` operates across columns and returns `x[:,1:size(x,1)]x[:, : size(x,1)-1]` for a 2-dimensional array. diff takes an optional keyword argument `axis` so that `diff(x, axis=0)` will operate across rows. diff can also be used to produce higher order differences (e.g. double difference).

```
>>> x= randn(3,4)
>>> x
array([[0.08542071,
 2.05598312,
 2.1114733 , 0.7986635 ],
 [0.17576066,
 0.83327885, 0.64064119,
 0.25631728],
 [0.38226593,
 1.09519101,
 0.29416551, 0.03059909]])
>>> diff(x) # Same as diff(x,1)
0.62339964288008698
>>> diff(x, axis=0)
array([[0.09033996,
 2.88926197, 2.75211449,
 1.05498078],
 [0.20650526,
 1.92846986,
 0.9348067 , 0.28691637]])
>>> diff(x, 2, axis=0) # Double difference, column by column
array([[0.11616531,
 4.81773183,
 3.68692119, 1.34189715]])
```

exp

exp returns the element-by-element exponential (e^x) for an array.

log

log returns the element-by-element natural logarithm ($\ln(x)$) for an array.

log10

log10 returns the element-by-element base-10 logarithm ($\log_{10}(x)$) for an array.

1.2 6.2 Rounding

around, round

around rounds to the nearest integer, or to a particular decimal place when called with two arguments.

```
%-----%
\begin{framed}
\begin{verbatim}
>>> x = randn(3)
array([ 0.60675173, 0.3361189
, 0.56688485])
>>> around(x)
array([ 1., 0., 1.])
>>> around(x, 2)
array([ 0.61, 0.34,
0.57])
\end{verbatim}
\end{framed}
```

around can also be used as a method on an ndarray – except that the method is named round. For example, `x.round(2)` is identical to `around(x, 2)`. The change of names is needed to avoid conflicting with the Python built-in function round.

floor

floor rounds to the next smallest integer.

```
>>> x = randn(3)
array([ 0.60675173, 0.3361189
```

```
, 0.56688485])  
>>> floor(x)  
array([ 0., 1.,  
1.])
```

ceil

ceil rounds to the next largest integer.

```
>>> x = randn(3)  
array([ 0.60675173, 0.3361189  
 , 0.56688485])  
>>> ceil(x)  
array([ 1., 0.,  
0.])
```

Note that the values returned are still floating points and so 0. is the same as 0..

1.3 Generating Arrays and Matrices

1.3.1 linspace

`linspace(l,u,n)` generates a set of `n` points uniformly spaced between `l`, a lower bound (inclusive) and `u`, an upper bound (inclusive).

```
>>> x = linspace(0, 10, 11)  
>>> x  
array([ 0., 1., 2., 3., 4., 5., 6., 7., 8., 9., 10.])
```

1.3.2 logspace

`logspace(l,u,n)` produces a set of logarithmically spaced points between 10^l and 10^u . It is identical to `10**linspace(l,u,n)`.

1.3.3 arange

`arange(l,u,s)` produces a set of points spaced by `s` between `l`, a lower bound (inclusive) and `u`, an upper bound (exclusive). `arange` can be used with a single parameter, so that

`arange(n)` is equivalent to `arange(0,n,1)`. Note that `arange` will return integer data type if all inputs are integer.

```
>>> x = arange(11)
array([ 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10])
>>> x = arange(11.0)
array([ 0., 1., 2., 3., 4., 5., 6., 7., 8., 9., 10.])
>>> x = arange(4, 10, 1.25)
array([ 4. , 5.25, 6.5 , 7.75, 9.  ])
```

1.3.4 meshgrid

`meshgrid` broadcasts two vectors to produce two 2-dimensional arrays, and is a useful function when plotting 3-dimensional functions.

```
>>> x = arange(5)
>>> y = arange(3)
>>> X,Y = meshgrid(x,y)
>>> X
array([[0, 1, 2, 3, 4],
       [0, 1, 2, 3, 4],
       [0, 1, 2, 3, 4]])
>>> Y
array([[0, 0, 0, 0, 0],
       [1, 1, 1, 1, 1],
       [2, 2, 2, 2, 2]])
```

1.3.5 r_

`r_` is a convenience function which generates 1-dimensional arrays from slice notation. While `r_` is highly flexible, the most common use is `r_[start : end : stepOrCount]` where `start` and `end` are the start and end points, and `stepOrCount` can be either a step size, if a real value, or a count, if complex.

```
>>> r_[0:10:1] # arange equiv
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

```
>>> r_[0:10:.5] # arange equiv
array([ 0. , 0.5, 1. , 1.5, 2. , 2.5, 3. , 3.5, 4. , 4.5, 5. ,
 5.5, 6. , 6.5, 7. , 7.5, 8. , 8.5, 9. , 9.5])
>>> r_[0:10:5j] # linspace equiv, includes end point
array([ 0. , 2.5, 5. , 7.5, 10. ])
```

`r_` can also be used to concatenate slices using commas to separate slice notation blocks.

```
>>> r_[0:2, 7:11, 1:4]
array([ 0, 1, 7, 8, 9, 10, 1, 2, 3])
```

Note that `r_` is not a function and that is used with `[]`.

1.3.6 `c_`

`c_` is virtually identical to `r_` except that column arrays are generated, which are 2-dimensional (second dimension has size 1)

```
>>> c_[0:5:2]
array([[0],
 [2],
 [4]])
>>> c_[1:5:4j]
array([[ 1. ],
 [ 2.33333333],
 [ 3.66666667],
 [ 5. ]])
```

`c_`, like `r_`, is not a function and is used with `[]`.

1.3.7 `ix_`

`ix_(a,b)` constructs an n -dimensional open mesh from n 1-dimensional lists or arrays. The output of `ix_` is an n -element tuple containing 1-dimensional arrays. The primary use of `ix_` is to simplify selecting slabs inside a matrix. Slicing can also be used to select elements from an array as long as the slice pattern is regular. `ix_` is particularly useful

for selecting elements from an array using indices which are not regularly spaced, as in the final example.

```
>>> x = reshape(arange(25.0),(5,5))
>>> x
array([[ 0.,  1.,  2.,  3.,  4.],
       [ 5.,  6.,  7.,  8.,  9.],
       [10., 11., 12., 13., 14.],
       [15., 16., 17., 18., 19.],
       [20., 21., 22., 23., 24.]])
>>> x[ix_([2,3],[0,1,2])] # Rows 2 & 3, cols 0, 1 and 2
array([[10., 11., 12.],
       [15., 16., 17.]])
>>> x[2:4,:3] # Same, standard slice
array([[10., 11., 12.],
       [15., 16., 17.]])
>>> x[ix_([0,3],[0,1,4])] # No slice equiv
```

1.3.8 mgrid

mgrid is very similar to meshgrid but behaves like `r_` and `c_` in that it takes slices as input, and uses a real valued variable to denote step size and complex to denote number of values. The output is an $n + 1$ dimensional vector where the first index of the output indexes the meshes.

```
>>> mgrid[0:3,0:2:.5]
array([[[ 0. ,  0. ,  0. ,  0. ],
       [ 1. ,  1. ,  1. ,  1. ],
       [ 2. ,  2. ,  2. ,  2. ]],
       [[ 0. ,  0.5,  1. ,  1.5],
       [ 0. ,  0.5,  1. ,  1.5],
       [ 0. ,  0.5,  1. ,  1.5]]])
>>> mgrid[0:3:3j,0:2:5j]
array([[[ 0. ,  0. ,  0. ,  0. ,  0. ],
       [ 1.5,  1.5,  1.5,  1.5,  1.5],
       [ 3. ,  3. ,  3. ,  3. ,  3. ]],
       [[ 0. ,  0.5,  1. ,  1.5,  2. ],
       [ 0. ,  0.5,  1. ,  1.5,  2. ],
       [ 0. ,  0.5,  1. ,  1.5,  2. ]]])
```

1.3.9 ogrid

ogrid is identical to mgrid except that the arrays returned are always 1-dimensional. ogrid output is generally more appropriate for looping code, while mgrid is usually more appropriate for vectorized code. When the size of the arrays is large, then ogrid uses much less memory.

```
>>> ogrid[0:3,0:2:.5]
[array([[ 0.],
        [ 1.],
        [ 2.]]) , array([[ 0. , 0.5, 1. , 1.5]])]
>>> ogrid[0:3:3j,0:2:5j]
[array([[ 0. ],
        [ 1.5],
        [ 3. ]]) ,
array([[ 0. , 0.5, 1. , 1.5, 2. ]])]
```