

1 Probability and Statistics Functions

1.1 Random Number Generation with NumPy

NumPy random number generators are all stored in the module `numpy.random`. These can be imported with using `import numpy as np` and then calling `np.random.rand`, for example, or by importing `import numpy.random as rnd` and using `rnd.rand.1`.

1.1.1 `rand`, `random_sample`

`rand` and `random_sample` are uniform random number generators which are identical except that `rand` takes a variable number of integer inputs – one for each dimension – while `random_sample` takes a n-element tuple.

`random_sample` is the preferred NumPy function, and `rand` is a convenience function primarily for MATLAB users.

```
>>> x = rand(3,4,5) >>> y = random_sample((3,4,5))
```

1.1.2 `randn`, `standard_normal`

`randn` and `standard_normal` are standard normal random number generators. `randn`, like `rand`, takes a variable number of integer inputs, and `standard_normal` takes an n-element tuple. Both can be called with no arguments to generate a single standard normal (e.g. `randn()`). `standard_normal` is the preferred NumPy function, and `randn` is a convenience function primarily for MATLAB users .

```
>>> x = randn(3,4,5)
>>> y = standard_normal((3,4,5))
```

1.1.3 `randint`, `random_integers`

`randint` and `random_integers` are uniform integer random number generators which take 3 inputs, `low`, `high` and `size`. `low` is the lower bound of the integers generated, `high` is the upper and `size` is a n-element tuple. `randint` and `random_integers` differ in that `randint` generates integers exclusive of the value in `high` (as do most Python functions), while `random_integers` includes the value in `high` in its range.

```
>>> x = randint(0,10,(100))
>>> x.max() # Is 9 since range is [0,10]
9
>>> y = random_integers(0,10,(100))
>>> y.max() # Is 10 since range is [0,10]
10
```

1.1.4 shuffle

shuffle randomly reorders the elements of an array in place.

```
>>> x = arange(10)
>>> shuffle(x)
>>> x
array([4, 6, 3, 7, 9, 0, 2, 1, 8, 5])
```

1.1.5 permutation

permutation returns randomly reordered elements of an array as a copy while not directly changing the input.

```
>>> x = arange(10)
>>> permutation(x)
array([2, 5, 3, 0, 6, 1, 9, 8, 4, 7])
>>> x
array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

NumPy provides a large selection of random number generators for specific distribution. All take between 0 and 2 required inputs which are parameters of the distribution, plus a tuple containing the size of the output. All random number generators are in the module `numpy.random`.

1.1.6 shuffle

shuffle randomly reorders the elements of an array in place.

```
>>> x = arange(10) >>> shuffle(x) >>> x array([4, 6, 3, 7, 9, 0, 2, 1, 8, 5])
```

1.1.7 permutation

`permutation` returns randomly reordered elements of an array as a copy while not directly changing the input.

```
>>> x = arange(10) >>> permutation(x) array([2, 5, 3, 0, 6, 1, 9, 8, 4, 7])
>>> x array([0, 1, 2, 3, 4, 5, 6, 7, 8, 9])
```

1.2 Select Random Number Generators

All take between 0 and 2 required inputs which are parameters of the distribution, plus a tuple containing the size of the output. All random number generators are in the module `numpy.random`.

1.2.1 `get_state`

`get_state()` gets the current state of the random number generator, which is a 5-element tuple. It can be called as a function, in which case it gets the state of the default random number generator, or as a method on a particular instance of `RandomState`.

1.2.2 `set_state`

`set_state(state)` sets the state of the random number generator. It can be called as a function, in which case it sets the state of the default random number generator, or as a method on a particular instance of `RandomState`. `set_state` should generally only be called using a state tuple returned by `get_state`.

1.3 19.2.2 Seed

`numpy.random.seed` is a more useful function for initializing the random number generator, and can be used in one of two ways. `seed()` will initialize (or reinitialize) the random number generator using some actual random data provided by the operating system.² `seed(s)` takes a vector of values (can be scalar) to initialize the random number generator at particular state. `seed(s)` is particularly useful for producing simulation studies which are reproducible. In the following example, calls to `seed()` produce different random numbers, since these reinitialize using random data from the computer, while calls to `seed(0)` produce the same (sequence) of random numbers.

```
>>> seed()
>>> randn()
array([ 0.62968838])
>>> seed()
>>> randn()
array([ 2.230155])
>>> seed(0)
>>> randn()
array([ 1.76405235])
>>> seed(0)
>>> randn()
array([ 1.76405235])
```

NumPy always calls `seed()` when the first random number is generated. As a result, calling `standard_normal()` across two “fresh” sessions will not produce the same random number.

1.4 19.3 Statistics Functions

1.4.1 mean

`mean` computes the average of an array. An optional second argument provides the axis to use (default is to use entire array). `mean` can be used either as a function or as a method on an array.

```
>>> x = arange(10.0)
>>> x.mean()
4.5
>>> mean(x)
4.5
>>> x= reshape(arange(20.0),(4,5))
>>> mean(x,0)
231
array([ 7.5,  8.5,  9.5, 10.5, 11.5])
>>> x.mean(1)
array([ 2.,  7., 12., 17.] )
```

1.4.2 median

`median` computed the median value in an array. An optional second argument provides the axis to use (default is to use entire array).

```
>>> x= randn(4,5)
>>> x
array([[0.74448693,
 0.63673031,
 0.40608815,
 0.40529852, 0.93803737],
 [ 0.77746525, 0.33487689, 0.78147524, 0.5050722
 , 0.58048329],
 [0.51451403,
 0.79600763,
 0.92590814, 0.53996231,
 0.24834136],
 [0.83610656,
 0.29678017, 0.66112691,
 0.10792584, 1.23180865]])
```

```
>>> median(x)
0.45558017286810903
>>> median(x, 0)
array([0.62950048,
0.16997507,
0.18769355, 0.19857318,
0.59318936])
```

Note that when an array or axis dimension contains an even number of elements (n), `median` returns the average of the 2 inner elements.

1.4.3 std

`std` computes the standard deviation of an array. An optional second argument provides the axis to use (default is to use entire array). `std` can be used either as a function or as a method on an array.

1.4.4 var

`var` computes the variance of an array. An optional second argument provides the axis to

....

1.4.5 corrcoef

`corrcoef(x)` computes the correlation between the rows of a 2-dimensional array `x`. `corrcoef(x, y)` computes the correlation between two 1-dimensional vectors. An optional keyword argument `rowvar` can be used to compute the correlation between the columns of the input – this is `corrcoef(x, rowvar=False)` and `corrcoef(x.T)` are identical

1.4.6 cov

`cov(x)` computes the covariance of an array `x`. `cov(x,y)` computes the covariance between two 1-dimensional vectors. An optional keyword argument `rowvar` can be used to compute the covariance between the columns of the input – this is `cov(x, rowvar=False)` and `cov(x.T)` are identical.

1.4.7 histogram

`histogram` can be used to compute the histogram (empirical frequency, using k bins) of a set of data. An optional second argument provides the number of bins. If omitted, $k = 10$ bins are used. `histogram` returns two outputs, the first with a k -element vector containing the number of observations in each bin, and the second with the $k + 1$ endpoints of the k bins.

1.4.8 normal

`normal()` generates draws from a standard Normal (Gaussian). `normal(mu, sigma)` generates draws from a Normal with mean μ and standard deviation σ . `normal(mu, sigma, (10,10))` generates a 10 by 10 array of draws from a Normal with mean μ and standard deviation σ .

`normal(mu, sigma)` is equivalent to `mu + sigma * standard_normal()`.

1.4.9 poisson

`poisson()` generates a draw from a Poisson distribution with $\lambda = 1$. `poisson(lambda)` generates a draw from a Poisson distribution with expectation λ . `poisson(lambda, (10,10))` generates a 10 by 10 array of draws from a Poisson distribution with expectation λ .

1.4.10 standard_t

`standard_t(nu)` generates a draw from a Student's t with shape parameter ν . `standard_t(nu, (10,10))` generates a 10 by 10 array of draws from a Student's t with shape parameter ν .

1.4.11 uniform

`uniform()` generates a uniform random variable on (0, 1). `uniform(low, high)` generates a uniform on (l , h). `uniform(low, high, (10,10))` generates a 10 by 10 array of uniforms on (l , h).

1.5 Simulation and Random Number Generation

Computer simulated random numbers are usually constructed from very complex but ultimately deterministic functions. Because they are not actually random, simulated random numbers are generally described to as **pseudo-random**.

All pseudo-random numbers in NumPy use one core random number generator based on the *Mersenne Twister*, a generator which can produce a very long series of pseudo-random data before repeating (up to $2^{19937} - 1$ non-repeating values) .

1.5.1 RandomState

RandomState is the class used to control the random number generators. Multiple generators can be initialized by RandomState.

```
>>> gen1 = np.random.RandomState()
>>> gen2 = np.random.RandomState()
>>> gen1.uniform() # Generate a uniform
0.6767614077579269
>>> state1 = gen1.get_state()
>>> gen1.uniform()
0.6046087317893271
>>> gen2.uniform() # Different, since gen2 has different seed
0.04519705909244154
>>> gen2.set_state(state1)
>>> gen2.uniform() # Same uniform as gen1 after assigning state
0.6046087317893271
```

1.5.2 19.2.1 State

Pseudo-random number generators track a set of values known as the state. The state is usually a vector which has the property that if two instances of the same pseudo-random number generator have the same state, the sequence of pseudo-random numbers generated will be identical. The state of the default random number generator can be read using `numpy.random.get_state` and can be restored using `numpy.random.set_state`.

```
>>> st = get_state()
>>> randn(4)
array([ 0.37283499, 0.63661908, 1.51588209,
        1.36540624])
>>> set_state(st)
```



```
>>> randn(4)
array([ 0.37283499, 0.63661908, 1.51588209,
        1.36540624])
```

The two sequences are identical since the state is the same when `randn` is called. The state is a 5- element tuple where the second element is a 625 by 1 vector of unsigned 32-bit integers. In practice the state should only be stored using `get_state` and restored using `set_state`.

1.6 normaltest

`normaltest` tests for normality in an array of data. An optional second argument provides the axis to use (default is to use entire array). Returns the test statistic and the p-value of the test. This test is a small sample modified version of the Jarque-Bera test statistic.

1.7 kstest

`kstest` implements the Kolmogorov-Smirnov test. Requires two inputs, the data to use in the test and the distribution, which can be a string or a frozen random variable object. If the distribution is provided as a string, then any required shape parameters are passed in the third argument using a tuple containing these parameters, in order.

```
>>> x = randn(100)
>>> kstest = stats.kstest
>>> stat, pval = kstest(x, 'norm')
>>> stat
0.11526423481470172
>>> pval
0.12963296757465059
>>> ncdf = stats.norm().cdf # No () on cdf to get the function
>>> kstest(x, ncdf)
(0.11526423481470172, 0.12963296757465059)
>>> x = gamma.rvs(2, size = 100)
>>> kstest(x, 'gamma', (2,)) # (2,) contains the shape parameter
(0.079237623453142447, 0.54096739528138205)
>>> gcdf = gamma(2).cdf
>>> kstest(x, gcdf)
(0.079237623453142447, 0.54096739528138205)
```

1.7.1 ks_2samp

`ks_2samp` implements a 2-sample version of the Kolmogorov-Smirnov test. It is called `ks_2samp(x,y)` where both inputs are 1-dimensional arrays, and returns the test statistic and p-value for the null that the distribution of `x` is the same as that of `y`. `shapiro` implements the Shapiro-Wilk test for normality on a 1-dimensional array of data. It returns the test statistic and p-value for the null of normality.

2 Statsmodels

`Statsmodels` is a Python module that allows users to explore data, estimate statistical models, and perform statistical tests. An extensive list of descriptive statistics, statistical tests, plotting functions, and result statistics are available for different types of data and each estimator. Researchers across fields may find that `statsmodels` fully meets their needs for statistical computing and data analysis in Python.

Features include:

- Linear regression models
- Generalized linear models
- Discrete choice models
- Robust linear models
- Many models and functions for time series analysis
- Nonparametric estimators
- A collection of datasets for examples
- A wide range of statistical tests
- Input-output tools for producing tables in a number of formats (Text, LaTeX, HTML) and for reading Stata files into NumPy and Pandas.
- Plotting functions
- Extensive unit tests to ensure correctness of results
- Many more models and extensions in development